

Report List-based sets

Author: Barau Elena-Raluca, Marculescu Tudor

Table of Contents

Introduction.....	3
Hand-over-Hand algorithm.....	4
Correctness of Hand-over-Hand algorithm.....	7
Performance analysis.....	9
Conclusions.....	14
References.....	15

Introduction

The goal of this project is to get familiar with the optimistic fine-grained locking schemes in concurrent data structures. Our running example is sorted linked list used to implement a set abstraction. The project was conducted in the context of “Fondements des algorithmes répartis” course. It proves that the performance is affected by the synchronization methods used on shared objects. Different ways of locking a set based list relates to the performance measured in a multi-threaded environment. Synchrobench, a java based software suite that test the performance of the locking schemes, was used to create the multi-threaded environment and to obtain the throughput as a performance metric. All the code and documents referenced in this report are present on the github repository from the references section. [1]

The algorithms investigated in this report are: CoarseGrainedListBasedSet, HandOverHandListIntSet and LazyLinkedListSortedSet. CoarseGrainedListBasedSet and LazyLinkedListSortedSet are already part of the Synchrobench suite, whereas HandOverHandListIntSet will be implemented and proved to be correct in this report.

In order to allow reproducibility of the results, both the hardware and the software environment are presented. Starting with the hardware environment that was used to conduct the experiments, namely a virtual machine from the infrastructure of the faculty, with the hostname **lame21.enst.fr**. It had the following specifications:

- Operating System: Ubuntu 24.04.3 LTS
Kernel: Linux 6.8.0-55-generic
- Processor: 2x Intel(R) Xeon(R) CPU E5-2660 v3 @ 2.60GHz
Architecture: x86-64
Per processor cores: 10 and threads: 2 per core, in total 20
Total cores 20, total threads 40
- Memory: 256 GiB DDR4
- Storage: HDD 256 GiB

The software environment software environment needed to run Synchrobench v1.1.0 is:

- GNU Make v4.3
- Apache Ant v1.9.5 [2]
- Java jdk v1.7.0 [3]
- Synchrobench release v1.1.0-alpha [4]

Please notice that Java 8 caused issues when trying to compile the project, particularly when the runtime jar was instrumented to enable transactional memory. We reverted to Java 7 as a fallback to avoid this issue.

Hand-over-Hand algorithm

The inspiration for implementing this algorithm came from the book “The art of multiprocessor programming” by Maurice Herlihy and Nir Shavit [5]. To avoid starting from scratch, the *CoarseGrainedListBasedSet* file was used as a reference and extensively modified to incorporate the required logic for the Hand-over-Hand algorithm implementation.

The methods that needed modifications are: *addInt(item)*, *removeInt(item)* and *containsInt(item)*. Each list class contains a Node class that defines how the node objects from the list are behaving. Previously, a lock was used in order to prevent the access of the other threads to the whole linked-list of Node objects, which is made obsolete by the Hand-over-Hand algorithm that uses a lock on each Node of the list. This is an algorithm with a fine grained approach on locking, therefore the Node class modified to accommodate this change. To showcase the implementation of the algorithm, it is sufficient to depict the methods of the novel class *HandOverHandListIntSet*, seen in the figures down below.

```
public boolean addInt(int item) {
    // try to lock the head
    head.lock();
    Node pred = head;
    // proceed to lock the next node
    try {
        Node curr = pred.next;
        // falls asleep until it can lock curr
        curr.lock();
        try {
            // traverse the list until finding the right position
            while (curr.key < item) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            // if the key is already present, return false
            if (curr.key == item) {
                return false;
            }
            // else, insert the new node
            Node newNode = new Node(item);
            newNode.next = curr;
            pred.next = newNode;
            return true;
        } finally {
            // unlock curr
            curr.unlock();
        }
    } finally {
        // unlock pred
        pred.unlock();
    }
}
```

Figure 1 – Hand-over-Hand algorithm insert method

```
public boolean removeInt(int item){
    Node pred = null, curr = null;

    // try to lock the head
    head.lock();
    try {
        pred = head;
        curr = pred.next;

        // try to lock curr
        curr.lock();
        try {
            // traverse the list until finding the right position
            while (curr.key < item) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            // if the key is present, remove it
            if (curr.key == item) {
                pred.next = curr.next;
                return true;
            }
            // if not found, return false
            return false;
        } finally {
            curr.unlock();
        }
    } finally {
        pred.unlock();
    }
}
```

Figure 2 – Hand-over-Hand algorithm remove method

For both add and remove methods, the list is traversed from head to tail by locking 2 successive nodes, the precedent and the current one. In this way, the algorithm allows for a boost in performance if some of the concurrent threads operate on nodes near the tail while others operate on those that are near the head. In contrast to the *CoarsedGrained* approach in which no matter where the operation of insert or remove took place, if one thread was executing either of them, the others had to busy-wait. By busy-waiting it is meant that the threads do not do any useful work, but waiting to acquire the lock on the linked-list.

The *lock()* and *unlock()* methods seen in the figures 1 and 2 are defined in the Node class. They are abstracting the calls to the *lock()* and *unlock()* functions of the *ReentrantLock* object. Each Node of *HandOverHandListIntSet* object contains a *ReentrantLock* object, which is seen in Figure 3.

```
private class Node {
    Node(int item) {
        key = item;
        next = null;
    }
    public int key;
    public Node next;
    public Lock lock = new ReentrantLock();

    public void lock() {
        lock.lock();
    }

    public void unlock() {
        lock.unlock();
    }
}
```

Figure 3 – Hand-over-Hand algorithm Node of the list

Finally, the *containsInt(item)* method is quite similar to the *removeInt(item)* method, except that it only searches for the specified item in the list without altering any of the nodes. The method is depicted in the Figure 4.

```
public boolean containsInt(int item){
    head.lock();
    Node pred = head;
    // proceed to lock the next node
    try {
        Node curr = pred.next;
        // falls asleep until it can lock curr
        curr.lock();
        try {
            // traverse the list until finding the right position
            while (curr.key < item) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            // if the key is already present, return false
            if (curr.key == item) {
                return true;
            }
            return false;
        } finally {
            curr.unlock();
        }
    } finally {
        pred.unlock();
    }
}
```

Figure 4 – Hand-over-Hand algorithm contains method

Correctness of Hand-over-Hand algorithm

In order to prove that the algorithm presented in the previous chapter is correct, it is sufficient to argue that both safety and liveness hold [6].

Starting with safety one can do the following assumptions: each Node has a lock that provides mutual exclusion. Acquiring a node's lock serializes access to that node and its next pointer; there are sentinel head ($-\text{MAX_INT}$) and tail ($+\text{MAX_INT}$) nodes so traversal never sees null and that the list is sorted strictly increasing and contains no duplicates; next pointers form a well-formed singly linked list.

A well-formed singly linked list refers to a properly constructed linked list in which: each node contains a valid data element and a reference (or pointer) to the next node; the list has a clear starting point (the head); the last node's next reference correctly points to null (or None), indicating the end of the list; there are no broken links (e.g., missing or invalid next references) and there are no cycles, meaning traversal from the head will eventually terminate.

Starting with the **addInt(item)** method, it is known that it traverses the list to locate pred and curr nodes such that $\text{pred.key} < \text{item} \leq \text{curr.key}$, with pred and curr locked when the check/insert is done. It has a check that prevents duplicates: if $\text{curr.key} == \text{item}$ return false while holding locks, which prevents a concurrent removal of the node.

One can argue that the **linearization point** for this method is the assignment $\text{pred.next} = \text{newNode}$, which is performed while pred and curr are locked. Before the write, $\text{pred.next} = \text{curr}$ and $\text{curr.key} \geq \text{item}$. After the write, newNode sits between pred and curr, with $\text{newNode.key} == \text{item}$ and $\text{newNode.next} = \text{curr}$, maintaining the list's sorted order and no duplicates, since the algorithm checks that $\text{curr.key} != \text{item}$. No concurrent writer could have changed pred.next or curr while we held pred and curr, so the **insertion is atomic at that moment**.

The **removeInt(item)** method locates pred and curr with $\text{pred.key} < \text{item} \leq \text{curr.key}$, with locks held. if curr.key is not item, it returns false.

For this operation, the **linearization point** is the assignment $\text{pred.next} = \text{curr.next}$, performed while pred and curr are locked. Removing curr splices it out while holding locks, so no concurrent writer can concurrently splice or insert at the same local place. The sorted order remains, since by removing an element from a sorted list, the list still remains sorted, and duplicates cannot be introduced, since this operation only removes items. If references to the removed node are kept elsewhere, memory reclamation must be handled safely, but in this case, it is the job of the garbage collector.

The **containsInt(item)** method traverses using lock coupling and tests curr.key . The linearization point is the test of $\text{curr.key} == \text{item}$ made while both pred and curr are locked. Because the node read is locked, no concurrent writer can change the relationship by inserting or removing in that local region between the read and the response. Thus the observed membership is consistent with some atomic instant in time.

Now taking all of the above arguments into consideration, one can prove that they produce a legal sequential history because: for any concurrent execution, by replacing each operation by its

linearization action at the linearization instant. Before that instant the operation had no visible effect; at that instant it performs an atomic write/read that corresponds to a legal sequential operation insert/remove/contains on a sorted set. After that it may release locks but further steps do not change the abstract set beyond that linearized effect. Therefore, each concurrent history is equivalent to some sequential history consistent with operation semantics. This proves that the concurrent histories of the algorithm are linearizable.

Liveness can be proved by arguing that the algorithm is deadlock free. This is a feature of the way locks are taken, from head to tail by each of the methods discussed for safety. They are enough to argue about since *addInt(item)*, *removeInt(item)* and *containsInt(item)* are the only operations executed concurrently at runtime. Figure 9.10, taken from the book “The art of multiprocessor programming” [5] depicts the case of a deadlock effect if the locks would have not been taken from head to tails in the case of all the methods.




Figure 9.10 The FineList class: a deadlock can occur if, for example, *remove()* and *add()* calls acquire locks in opposite order. Thread A is about to insert *a* by locking first *b* and then *head*, and thread B is about to remove node *b* by locking first *head* and then *b*. Each thread holds the lock the other is waiting to acquire, so neither makes progress.

Performance analysis

For the performance analysis, a graph was created showing the throughput as a function of the number of threads for all three algorithms: CoarseGrainedListBasedSet, HandOverHandListIntSet and LazyLinkedListSortedSet, using representative values for list size and update ratio. Then, for each algorithm, the update ratio was fixed at 10%, and a graph was plotted showing the throughput versus the number of threads while varying the list size. Finally, for each algorithm, a graph was prepared showing the throughput versus the number of threads while varying the update ratio, using a list size of 1k.

In order to measure the performance of the algorithms, the following command was used to call the synchrobench algorithm and obtain the throughput in **operations / second**.

```
java -cp bin contention.benchmark.Test -b linkedlists.lockbased.$OUTPUT -d 2000  
-t $i -u $j -i $k -r $r -W 0 | grep Throughput
```

where:

```
$OUTPUT == name of the class containing the algorithm to be used  
-t $i    == sets the number of threads to be used while doing the experiment  
-u $j    == sets the update ratio in percentage while doing the experiment  
-i $k    == sets the initial list size  
-r $r    == sets the maximal list size, the range  
-W 0     == 0 seconds for warmup  
-d 2000  == 2s duration of each experiment
```

In the following figures, the plots are representative for each algorithm with a fixed update ratio of 10% but with varying list sizes.




Figure 5: Coarse-Grained Locking algorithm for a fixed update ratio 10% and varying list size




Figure 6: Hand-Over-Hand Locking algorithm for a fixed update ratio 10% and varying list size




Figure 7: Lazy Linked List algorithm for a fixed update ratio 10% and varying list size

The following plots depict each algorithm with a fixed list size but varying update ratios.




Figure 8: Coarse-Grained Locking algorithm for a fixed list size 100 and varying update ratios

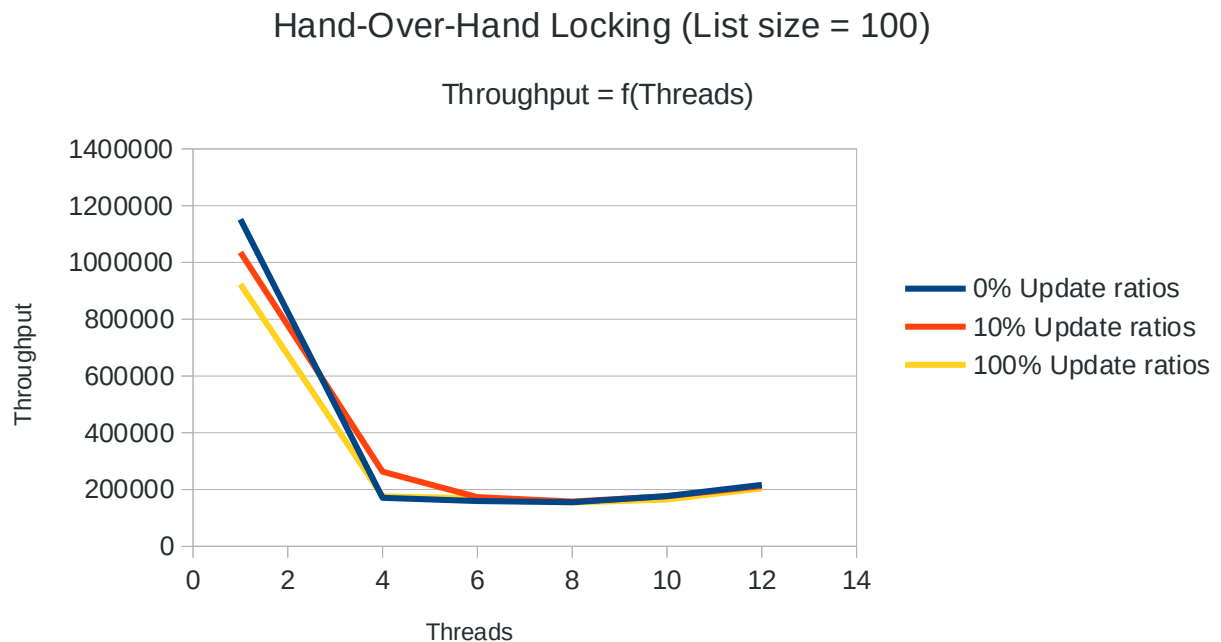


Figure 9: Hand-Over-Hand Locking algorithm for a fixed list size 100 and varying update ratios




Figure 10: Lazy Linked List algorithm for a fixed list size 100 and varying update ratios
Finally one plot, with three curves, one for each algorithm, with fixed update ratio 10% and list size 1000.




Figure 11: Throughput as the function of the number of threads for all three algorithms, with fixed update ratio 10% and list size 1000

// Explain the forms of the curves and their relations to each other.

// Shall be done after check of the corectness

Conclusions

In this project, we explored the implementation and analysis of list-based sets using different synchronization strategies in a concurrent environment. By focusing on the Hand-over-Hand (fine-grained) locking algorithm, we demonstrated how carefully managing locks at the node level can significantly reduce contention compared to coarse-grained locking.

The correctness of the Hand-over-Hand algorithm was established through linearizability arguments, showing that all operations maintain the invariants of a sorted, duplicate-free list while preventing deadlocks through ordered lock acquisition. Safety was ensured by holding locks during critical updates, and liveness considerations highlighted the trade-offs between blocking progress and throughput under high contention.

Performance analysis confirmed that fine-grained locking provides better scalability as the number of threads increases, especially for operations distributed across different parts of the list. However, some limitations remain, such as potential contention near the head of the list and the need for fair locks or optimistic approaches to improve throughput further.

Overall, this study illustrates that synchronization strategies have a direct and measurable impact on the efficiency of concurrent data structures. Fine-grained approaches like Hand-over-Hand locking offer a practical balance between correctness, safety, and performance, making them suitable for multi-threaded applications requiring shared access to linked-list-based sets.

References

1. *optimistic-lock-based-list-based-set*, <https://github.com/marcutudor79/optimistic-lock-based-list-based-set> (last checked: 27.10.2025)
2. Apache, *The Apache ANT Project*, v.1.9.4. <https://ant.apache.org/bindownload.cgi> (last checked: 27.10.2025)
3. Oracle, *Java SE Development Kit 7u80*, v.1.7.0. <https://www.oracle.com/in/java/technologies/javase/javase7-archive-downloads.html> (last checked: 27.10.2025)
4. Synchrobench, *A Java-based benchmarking suite for concurrent data structures*, v1.1.0-alpha. <https://github.com/gramoli/synchrobench> (last checked: 27.10.2025)
5. M. Herlihy and N. Shavit, *The Art of Multiprocessor Programming*, USA: Morgan Kaufmann, 2008
6. P. Kuznetsov, *Concurrent List-Based Sets*, CSC_4SL05_TP, 2020