

Report List-based sets

Author: Barau Elena-Raluca, Marculescu Tudor

Table of Contents

Introduction.....	3
Hand-over-Hand algorithm.....	4
Correctness of Hand-over-Hand algorithm.....	5
Performance analysis.....	6
Conclusions.....	7
References.....	8

Introduction

The goal of this project is to get familiar with optimistic fine-grained locking schemes in concurrent data structures. Our running example is sorted linked list used to implement a set abstraction. The project was conducted in the context of “Fondements des algorithmes répartis (partie A)” course.

It proves that the performance is affected by the synchronization methods used on shared objects. The purpose of this project is to show how different ways of locking a set based list relates to the performance measured in a multi-threaded environment.

The hardware environment used was a virtual machine from the infrastructure of the faculty, namely: **lame21.enst.fr** which has the following specifications:

- Operating System: Ubuntu 24.04.3 LTS
Kernel: Linux 6.8.0-55-generic
- Processor: 2x Intel(R) Xeon(R) CPU E5-2660 v3 @ 2.60GHz
Architecture: x86-64
Per processor cores: 10 and threads: 2 per core, in total 20
Total cores 20, total threads 40
- Memory: 256 GiB
- Storage: HDD 256 GiB

The software environment used to run synchrobench, which is the java based software suite to test the performance of our shared objects is:

- GNU Make v4.3
- Apache Ant v1.9.5
- Java jdk v1.7.0
- synchrobench release v1.1.0-alpha

Please notice that Java 8 caused issues when compiling the synchrobench project, particularly when the runtime jar was instrumented to enable transactional memory. We resorted to fallback to Java 7 in order to prevent this issue.

Hand-over-Hand algorithm

The inspiration of implementing this algorithm came from the book “The art of multiprocessor programming” by Maurice Herlihy and Nir Shavit. In order to not start from scratch, the CoarseGrainedListBasedSet file was taken as an example and have been heavily modified in order to introduce the necessary logic of implementation for Hand-over-Hand algorithm.

The methods that have been modified are: addInt(item), removeInt(item) and containsInt(item). Previously, a lock was used in order to prevent the access of the other threads to the whole linked-list of Node objects, which is made obsolete by the Hand-over-Hand algorithm that uses a lock on each Node of the list, thus being an algorithm with a fine grained approach on locking. The Node class was also modified to accommodate this change. To showcase the implementation of the algorithm, it is sufficient to depict the methods of the novel class HandOverHandListIntSet as seen in the figures down below.

```
public boolean addInt(int item) {
    // try to lock the head
    head.lock();
    Node pred = head;
    // proceed to lock the next node
    try {
        Node curr = pred.next;
        // falls asleep until it can lock curr
        curr.lock();
        try {
            // traverse the list until finding the right position
            while (curr.key < item) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            // if the key is already present, return false
            if (curr.key == item) {
                return false;
            }
            // else, insert the new node
            Node newNode = new Node(item);
            newNode.next = curr;
            pred.next = newNode;
            return true;
        } finally {
            // unlock curr
            curr.unlock();
        }
    } finally {
        // unlock pred
        pred.unlock();
    }
}
```

Figure 1 – Hand-over-Hand algorithm insert method

```
public boolean removeInt(int item){
    Node pred = null, curr = null;

    // try to lock the head
    head.lock();
    try {
        pred = head;
        curr = pred.next;

        // try to lock curr
        curr.lock();
        try {
            // traverse the list until finding the right position
            while (curr.key < item) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            // if the key is present, remove it
            if (curr.key == item) {
                pred.next = curr.next;
                return true;
            }
            // if not found, return false
            return false;
        } finally {
            curr.unlock();
        }
    } finally {
        pred.unlock();
    }
}
```

Figure 2 – Hand-over-Hand algorithm remove method

For both methods, the list is traversed from head to tail by locking 2 successive nodes, the precedent and the current one. In this way, the algorithm allows for a boost in performance if some of the concurrent threads operate on nodes near the tail while others operate on nodes that are near the head, compared to the CoarsedGrained approach in which no matter where the operation of insert or remove took place, if one thread was executing either of them, the others had to busy-wait. By busy-waiting it is meant that the threads do not do any useful work, but waiting to acquire the lock on the linked-list.

The lock() and unlock() methods seen in the figures 1 and 2 are defined in the Node class, that is part of the HandOverHandListIntSet class. They are abstracting the calls to the lock() and unlock() functions of the ReentrantLock object. Each Node object contains a ReentrantLock object in the case of HandOverHandListIntSet as seen in Figure 3.

```
private class Node {
    Node(int item) {
        key = item;
        next = null;
    }
    public int key;
    public Node next;
    public Lock lock = new ReentrantLock();

    public void lock() {
        lock.lock();
    }

    public void unlock() {
        lock.unlock();
    }
}
```

Figure 3 – Hand-over-Hand algorithm Node of the list

It remains to showcase one final relevant method for synchronization, which is the `containsInt(item)` method. It is not much different than the `removeInt(item)` method, but it only searches for the given item in the list, without modifying the Nodes. The method is depicted in the Figure 4.

```
public boolean containsInt(int item){
    head.lock();
    Node pred = head;
    // proceed to lock the next node
    try {
        Node curr = pred.next;
        // falls asleep until it can lock curr
        curr.lock();
        try {
            // traverse the list until finding the right position
            while (curr.key < item) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            // if the key is already present, return false
            if (curr.key == item) {
                return true;
            }
            return false;
        } finally {
            // unlock curr
            curr.unlock();
        }
    } finally {
        // unlock pred
        pred.unlock();
    }
}
```

Figure 4 – Hand-over-Hand algorithm contains method

Correctness of Hand-over-Hand algorithm

[] Argue that both safety and liveness hold

Performance analysis

For the performance analysis, prepare a graph depicting the throughput as a function of the number of threads for the three algorithms, for some representative list size and update ratio. You may use gnuplot or any other plotting program of your choice. Then, for each algorithm, fix the update ratio to 10%, and prepare a graph depicting the throughput as the function of the number of threads, varying the list size. Finally, for each algorithm, prepare a graph depicting the throughput as the function of the number of threads, varying the update ratio, for the list size 1k. Altogether, this gives:

- [] Three plots (one per algorithm), with three curves each, for a fixed update ratio 10% and varying list size.
- [] Three plots (one per algorithm), with three curves each, for a fixed list size 100 and varying update ratios.
- [] One plot, with three curves (one per algorithm), with fixed update ratio 10% and list size 1000.
- [] Explain the forms of the curves and their relations to each other.

Conclusions

References